

TITLE : To perform Exploratory Data Analysis (EDA), enhance classification accuracy using ensemble

methods, and evaluate model performance using appropriate validation techniques in Python.

NAME : Shinde Shubham Dnyandev

ROLL NO. : 23107121

CLASS : TY-B

BATCH : B

In [2]:

```
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.datasets import load_breast_cancer
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier, VotingClassifier
from sklearn.metrics import classification_report, confusion_matrix, accuracy_score,
from sklearn.preprocessing import StandardScaler
```

In [4]:

```
cancer = load_breast_cancer()
df = pd.DataFrame(cancer.data, columns=cancer.feature_names)
df['target'] = cancer.target

print("\nDataset shape:", df.shape)
print("\nFirst 5 rows:\n", df.head())
print("\nClass distribution:\n", df['target'].value_counts())
```

```
Dataset shape: (569, 31)
```

```
First 5 rows:
```

```
    mean radius  mean texture  mean perimeter  mean area  mean smoothness \
0      17.99      10.38       122.80     1001.0      0.11840
1      20.57      17.77       132.90     1326.0      0.08474
2      19.69      21.25       130.00     1203.0      0.10960
3      11.42      20.38        77.58      386.1      0.14250
4      20.29      14.34       135.10     1297.0      0.10030
```

```
    mean compactness  mean concavity  mean concave points  mean symmetry \
0      0.27760      0.3001       0.14710      0.2419
1      0.07864      0.0869       0.07017      0.1812
2      0.15990      0.1974       0.12790      0.2069
3      0.28390      0.2414       0.10520      0.2597
4      0.13280      0.1980       0.10430      0.1809
```

```
    mean fractal dimension  ...  worst texture  worst perimeter  worst area \
0      0.07871     ...      17.33       184.60     2019.0
1      0.05667     ...      23.41       158.80     1956.0
2      0.05999     ...      25.53       152.50     1709.0
3      0.09744     ...      26.50        98.87     567.7
4      0.05883     ...      16.67       152.20     1575.0
```

```
    worst smoothness  worst compactness  worst concavity  worst concave points \
0      0.1622       0.6656       0.7119      0.2654
1      0.1238       0.1866       0.2416      0.1860
2      0.1444       0.4245       0.4504      0.2430
3      0.2098       0.8663       0.6869      0.2575
4      0.1374       0.2050       0.4000      0.1625
```

```
    worst symmetry  worst fractal dimension  target
0      0.4601       0.11890      0
1      0.2750       0.08902      0
2      0.3613       0.08758      0
3      0.6638       0.17300      0
4      0.2364       0.07678      0
```

```
[5 rows x 31 columns]
```

```
Class distribution:
```

```
target
1    357
0    212
Name: count, dtype: int64
```

```
In [6]: X = df.drop('target', axis=1)
y = df['target']
```

```
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
```

```
In [8]: X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2, random_state=42)
```

```
In [10]: rf = RandomForestClassifier(n_estimators=100, random_state=42)
gb = GradientBoostingClassifier(n_estimators=100, random_state=42)
ensemble = VotingClassifier(estimators=[('rf', rf), ('gb', gb)], voting='soft')

models = {'Random Forest': rf, 'Gradient Boosting': gb, 'Voting Ensemble': ensemble}
```

```
In [12]: for name, model in models.items():
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)

    print(f"\n==== {name} ====")
```

```
print("Accuracy:", accuracy_score(y_test, y_pred))
print("ROC-AUC:", roc_auc_score(y_test, model.predict_proba(X_test)[:, 1]))
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred))
print("Classification Report:\n", classification_report(y_test, y_pred))
```

==== Random Forest ===

Accuracy: 0.956140350877193

ROC-AUC: 0.9937169312169312

Confusion Matrix:

```
[[39  3]
 [ 2 70]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.95	0.93	0.94	42
1	0.96	0.97	0.97	72
accuracy			0.96	114
macro avg	0.96	0.95	0.95	114
weighted avg	0.96	0.96	0.96	114

==== Gradient Boosting ===

Accuracy: 0.956140350877193

ROC-AUC: 0.9907407407407407

Confusion Matrix:

```
[[38  4]
 [ 1 71]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.97	0.90	0.94	42
1	0.95	0.99	0.97	72
accuracy			0.96	114
macro avg	0.96	0.95	0.95	114
weighted avg	0.96	0.96	0.96	114

==== Voting Ensemble ===

Accuracy: 0.956140350877193

ROC-AUC: 0.9927248677248678

Confusion Matrix:

```
[[38  4]
 [ 1 71]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.97	0.90	0.94	42
1	0.95	0.99	0.97	72
accuracy			0.96	114
macro avg	0.96	0.95	0.95	114
weighted avg	0.96	0.96	0.96	114

In [13]:

```
# Cross-validation accuracy
cv_score = cross_val_score(model, X_scaled, y, cv=5, scoring='accuracy').mean()
print("Cross-validation Accuracy:", cv_score)
```

Cross-validation Accuracy: 0.9613569321533924

In [15]:

```
feat_importances = pd.Series(rf.feature_importances_, index=cancer.feature_names)
top_features = feat_importances.nlargest(10).index

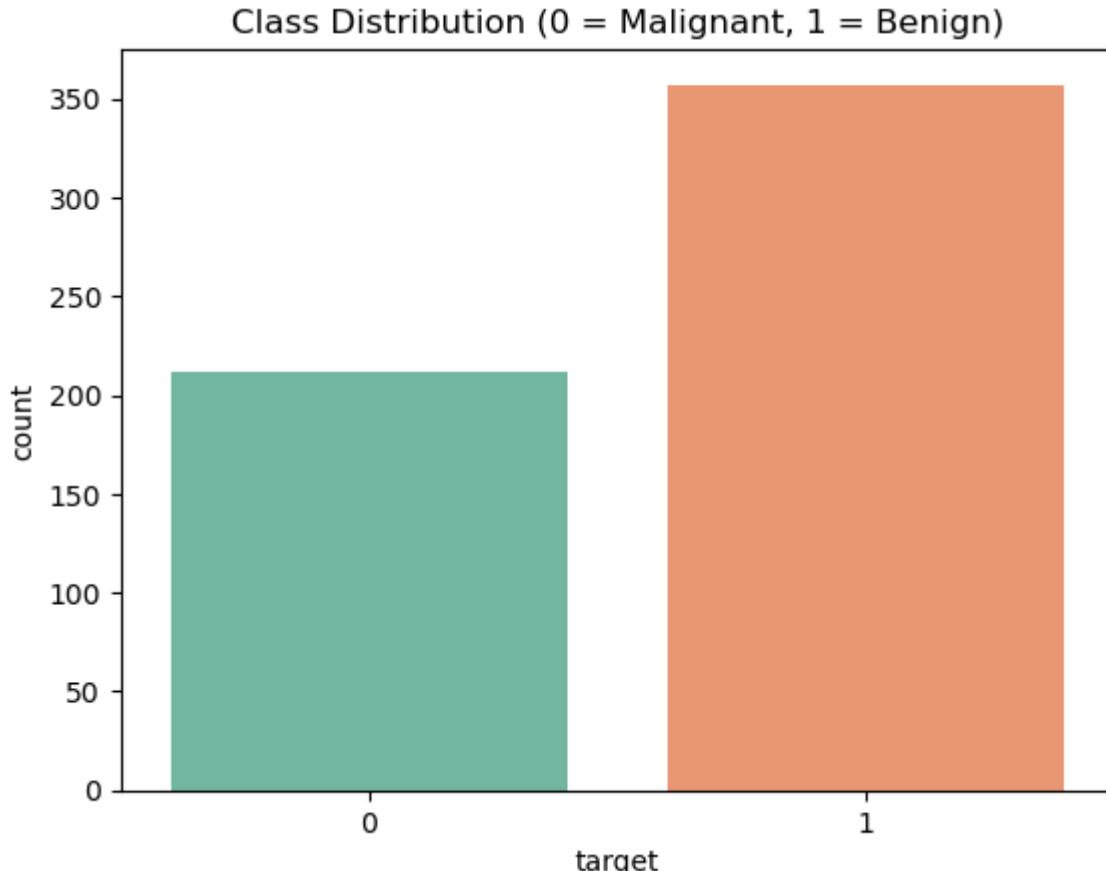
sns.countplot(x='target', data=df, palette='Set2')
```

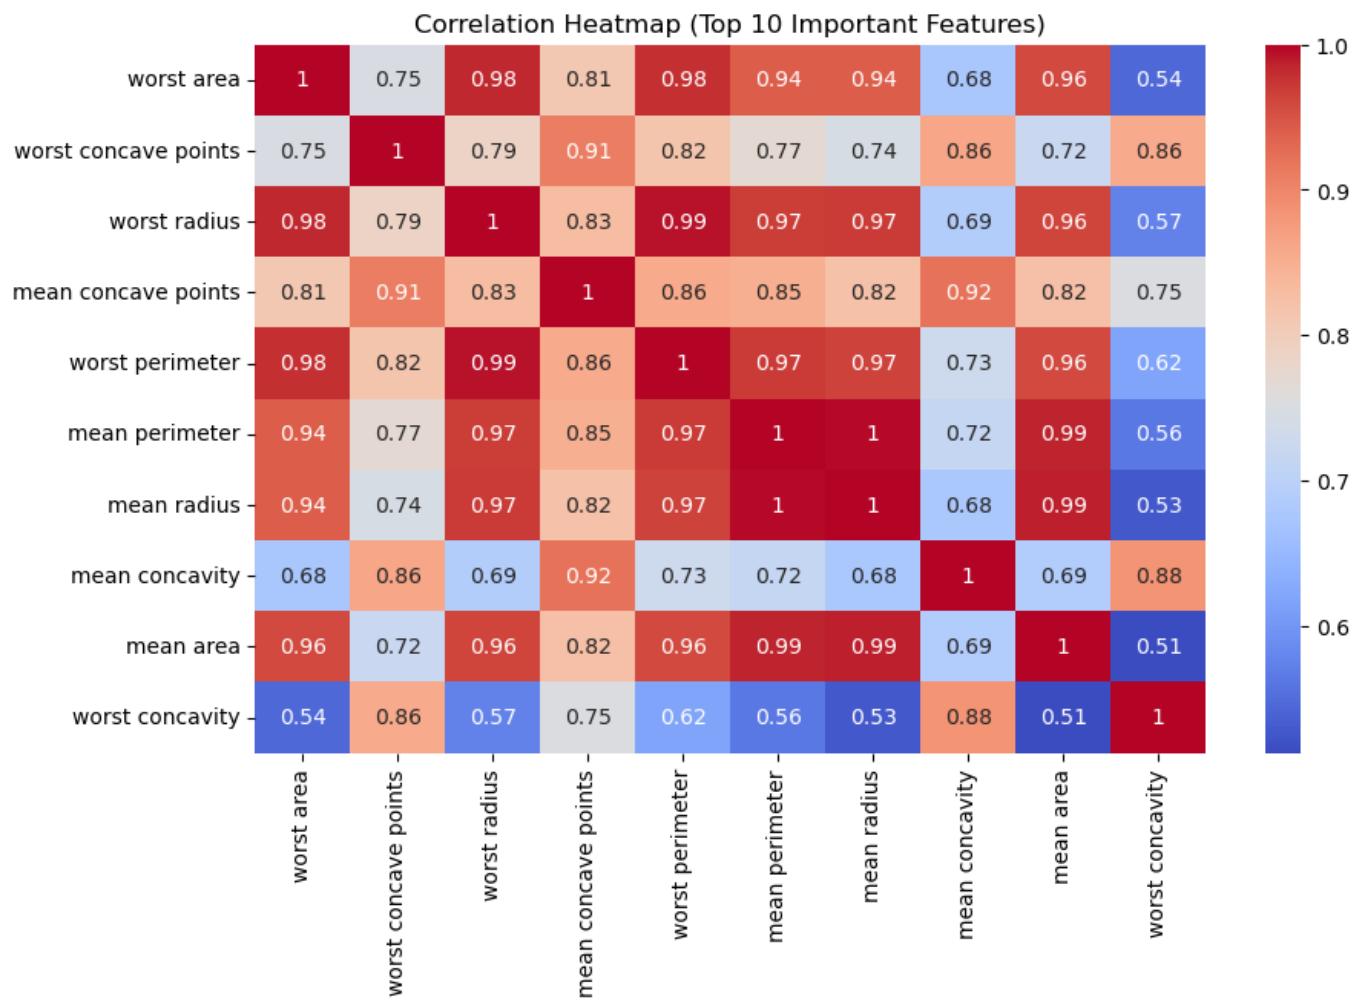
```
plt.title("Class Distribution (0 = Malignant, 1 = Benign)")  
plt.show()  
  
plt.figure(figsize=(10, 6))  
sns.heatmap(df[top_features].corr(), annot=True, cmap="coolwarm", cbar=True)  
plt.title("Correlation Heatmap (Top 10 Important Features)")  
plt.show()
```

/tmp/ipykernel_3111/1012119997.py:4: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.countplot(x='target', data=df, palette='Set2')
```





In []: